

# Introduction to Amazon DynamoDB

- and some use case examples

Pete Naylor  
Sr Database Specialist SA - AWS

# Agenda


- DynamoDB's place in the history of purpose-built databases
- Key concepts
- How do I use DynamoDB?
  - Basic data models and controls
- Challenges with changing and imbalanced loads
  - And how DynamoDB deals with them
- Integrated DynamoDB architecture and data flows

# Database evolution and DynamoDB

# DynamoDB history

## How did we get here?

Dyna



Andy Jassy

@ajassy

Following


Store

We migrated 75 petabytes of internal data stored in nearly 7,500 Oracle databases to multiple AWS database services including [Amazon DynamoDB](#), [Amazon Aurora](#), [Amazon Relational Database Service \(RDS\)](#), and [Amazon Redshift](#). The migrations were accomplished with little or no downtime, and covered 100% of our proprietary systems. This includes complex purchasing, catalog management, order fulfillment, accounting, and video streaming workloads. We kept careful track of the costs and the performance, and realized the following results:

- **Cost Reduction** – We reduced our database costs by over 60% on top of the heavily discounted rate we negotiated based on our scale. Customers regularly report cost savings of 90% by switching from Oracle to AWS.
- **Performance Improvements** – Latency of our consumer-facing applications was reduced by 40%.
- **Administrative Overhead** – The switch to managed services reduced database admin overhead by 70%.

around the world. A continuously and the of these failures d software systems.

868 Retweets 2,281 Likes




76

868

2.3K

content are being ce responsible for always write to and eds to be available

© 2020, Amazon Web Services, Inc. or its Affiliates.



# DynamoDB history

## Motivations at Amazon (and elsewhere)

1. Reduce dependence on commercially licensed relational database engines
2. Minimize operational complexity and administrative overhead
3. Provide best possible customer experience across all data indexing needs

“A one size fits all database doesn't fit anyone”

- Werner Vogels

# Internet-scale applications


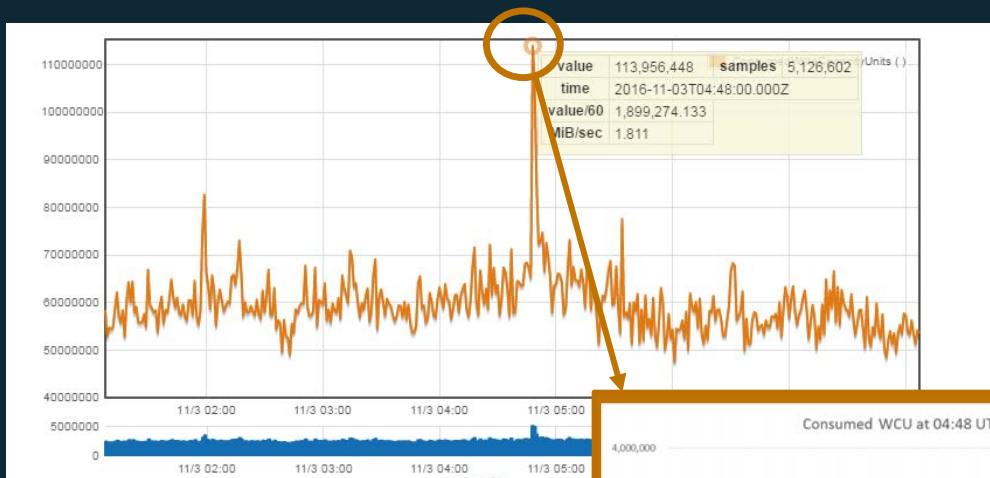


Users	1M+
Data volume	TB-PB-EB
Locality	Global
Performance	Milliseconds-microseconds
Request rate	Millions
Access	Mobile, IoT, devices
Scale	Incrementally based on traffic
Economics	Pay as you go
Developer access	Instant API access






# Snap (Snapchat)



Database writes peak *seconds* after Chicago Cubs win the World Series.



# A migration from mainframe



- Retail business ran on mainframe, which became a bottleneck
- Migrated financial transaction data to DynamoDB
- Unbound scale for customers and app developers

“We built a secure and resilient cloud infrastructure that could solve the scalability and reliability problems with a serverless architecture.”

Srini Uppalapati  
Capital One

# SQL and DynamoDB side by side

## Traditional SQL

## DynamoDB

Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad hoc queries	Instantiated views for known patterns
Scale vertically	Scale horizontally
Good for OLAP	Built for OLTP at scale

# Core Strengths of DynamoDB

## Sounds cool, but when should I use it?

### Priorities

Security

Durability

Scalability

Availability

Low Latency

Easy to Use

Easy to Manage

Considerations:

- SQL / NoSQL
- Relational / Non-Relational
- Operational / Analytical


DynamoDB solves for:

- Horizontal scaling
- Decoupled compute/storage
- Asynchronous roll-ups/aggregations
- Availability/Durability/Latency
- Well-known access patterns
- Schema flexibility

# Key concepts


# Scaling databases

## Traditional SQL



Scale up


## NoSQL



Scale out to many shards


# Scaling NoSQL databases

## Most NoSQL databases



Servers and clusters

## DynamoDB



DynamoDB: partitions

Basic premise: There is a way to shard data that's horizontally scalable.






# Incremental scaling with DynamoDB

Workload:  
data volume, reads, writes


DynamoDB resources:  
storage, read, and write capacity




# You work with tables...



## DynamoDB does the rest under the hood...

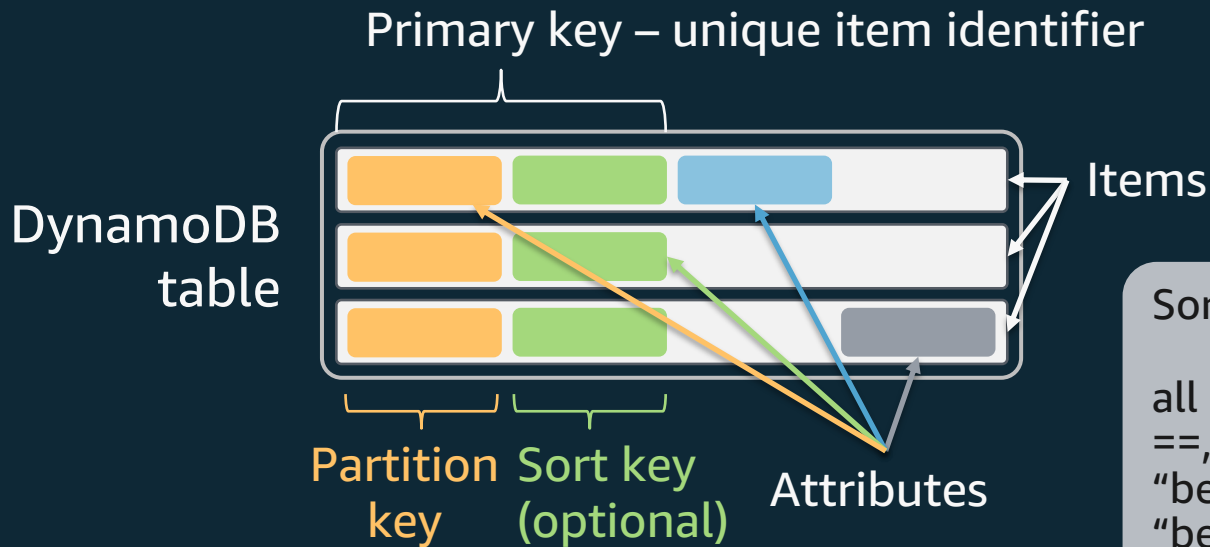


...



1K WU or 3K RU  
up to 10 GB

# Terminology



Sort key conditions:

all items  
==, <, >, >=, <=  
"begins with"  
"between"

sorted results  
counts  
top/bottom N

- 1:1 relationships
- distribute traffic
- collection identity

- 1:many relationships
- collect related items
- efficient filtering
- sorting

# Global secondary index (GSI)


Alternate partition (+sort) key for alternate materialized view

Online indexing

Capacity provisioned separately from table

Up to 20 GSIs per table

Table



GSIs



*KEYS\_ONLY*



*INCLUDE A3*

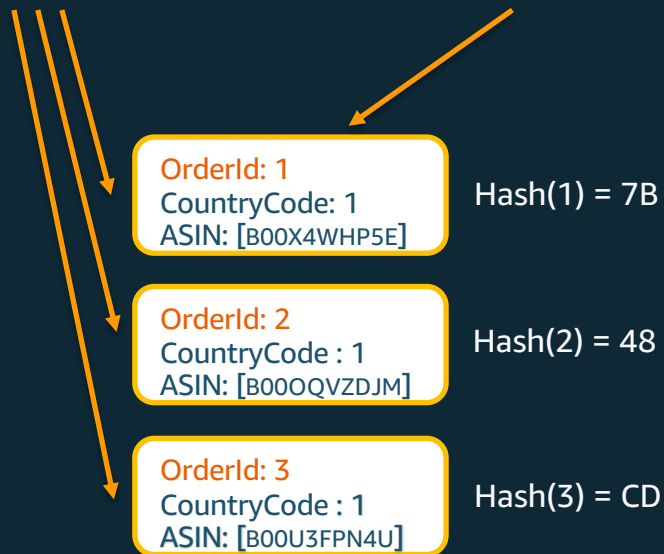


*ALL*

# Sharding/partitioning

Denormalized Record

Partition key

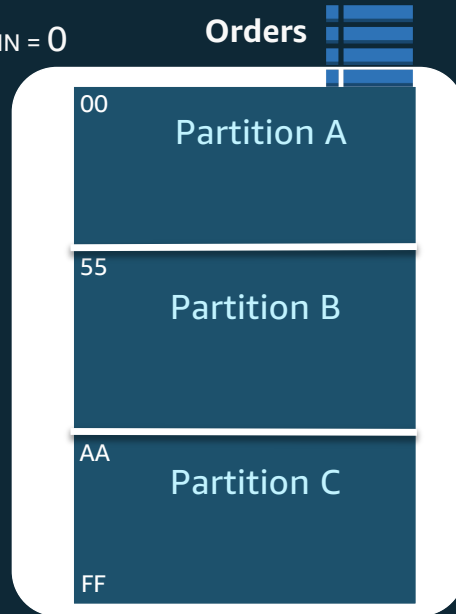


Keyspace

DynamoDB table

Hash.MIN = 0

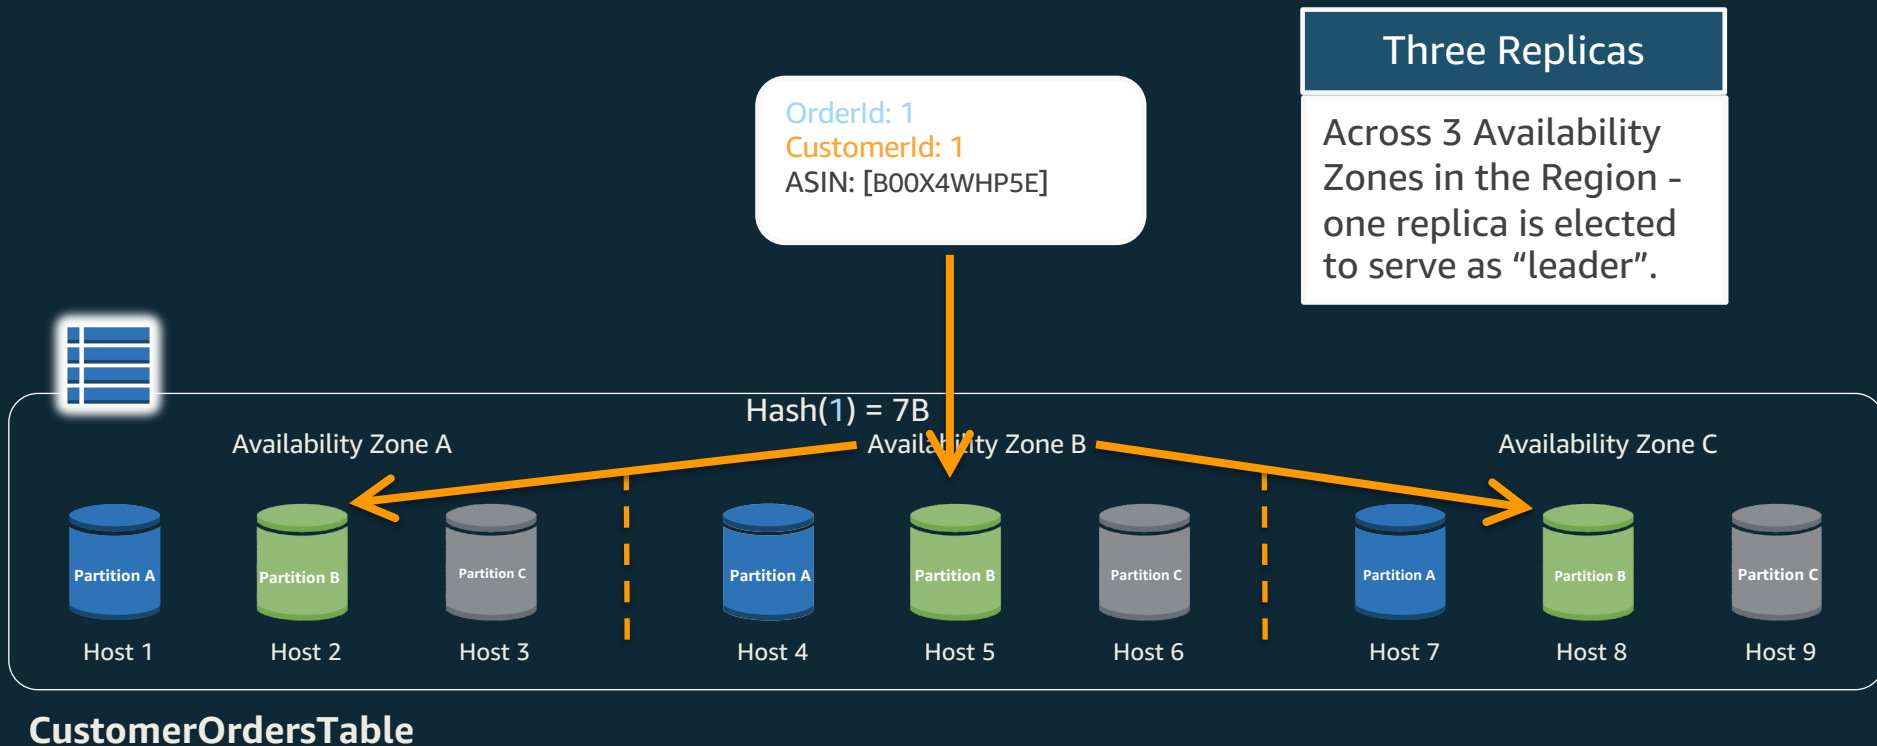
Orders



Hash.MAX = FF

Related data is stored together for efficient access

# A view “from a different angle”





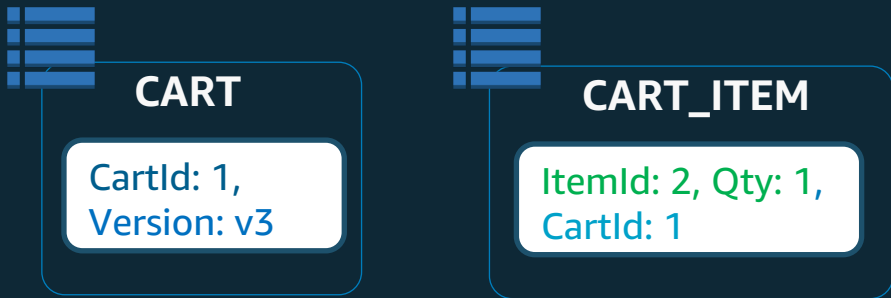
# Modeling data for DynamoDB

# Partition key

- A good sharding (partitioning) scheme affords even distribution of both data and workload as they grow
- Key concept: **partition key** as the dimension of scalability
  - **Distribute traffic and data** across partitions – horizontal scaling
- Ideal scaling conditions:
  - The partition key is from a high cardinality set (that grows)
  - Requests are evenly spread over the key space
  - Requests are evenly spread over time

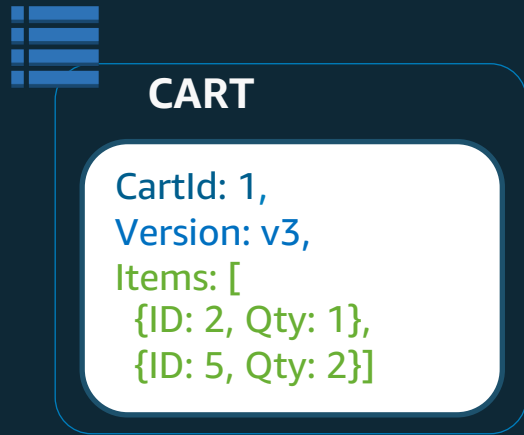
# Denormalization

## Traditional database



Normalized schema: multiple items  
– one table per entity type

## DynamoDB



Denormalized: one item  
with flexible schema

# Ensuring data consistency on updates

## Example: add/remove shopping cart items

1. get cart =>  $v_{\text{read}} = \text{Version}$

2. update cart:

**IF**  $\text{Version} = v_{\text{read}}$   
    add/remove cart items  
    ++Version

**ELSE** go back to Step 1.

Use **ConditionExpression**  
in DynamoDB

Optimistic concurrency control

# Updating cart: data consistency using OCC

## 1. Get the cart: GetItem

```
{  "TableName": "Cart",
  "Key": {"CartId": {"N": "2"}}
}
```



**CART**

```
CartID:2,
Version: 3,
CartItems: [
  {ID: 2, Qty: 1},
  {ID: 5, Qty: 2}]
```

## 2. Update the cart: conditional PutItem (or UpdateItem)

```
{  "TableName": "Cart",
  "Item": {
    "CartID": {"N": "2"},
    "Version": {"N": "4"},
    "CartItems": {...}
  },
  "ConditionExpression": "Version = :ver",
  "ExpressionAttributeValues": {":ver": {"N": "3"}}
}
```

- ✓ Use conditions to implement optimistic concurrency control ensuring data consistency
- ✓ Single-item operations are ACID
- ✓ GetItem call can be eventually consistent

Yes, you can have strongly consistent  
read-after-write and concurrency control  
with DynamoDB



# One-to-one relationships or key-values

- Use a table or GSI with a partition key
- Use GetItem or BatchGetItem API

Example: Given a user or email, get attributes

Users table	
Partition key	Attributes
UserId = bob	Email = bob@example.com, JoinDate = 2011-11-15
UserId = fred	Email = fred@example.com, JoinDate = 2011-12-01

Users-Email-GSI	
Partition key	Attributes
Email = <u>bob@example.com</u>	UserId = bob, JoinDate = 2011-11-15
Email = <u>fred@example.com</u>	UserId = fred, JoinDate = 2011-12-01

# One-to-many relationships or parent-children

- Use a table or GSI with a partition and sort key
- Use the Query API to get multiple items

Example: Given a device, find all readings between epoch X, Y

Device-measurements		
Part. key	Sort key	Attributes
DeviceId = 1	epoch = 5513A97C	Temperature = 30, pressure = 90
DeviceId = 1	epoch = 5513A9DB	Temperature = 30, pressure = 90

# Many-to-many relationships

- Use a table and GSI with the partition and sort key elements switched
- Use the Query API

Example: Given a user, find all games. Or given a game, find all users.


User-Games-Table	
Part. key	Sort key
UserId = bob	GameId = Game1
UserId = fred	GameId = Game2
UserId = bob	GameId = Game3

Game-Users-GSI	
Part. key	Sort key
GameId = Game1	UserId = bob
GameId = Game2	UserId = fred
GameId = Game3	UserId = bob

Yes, you can model complex data  
relationships with DynamoDB

# Challenges with growing datasets, variable throughput, imbalanced workloads


# Time-To-Live (TTL)





# Provisioned Mode with Auto Scaling


Reads




Writes



# On-demand Mode: Spiky workloads




# Imbalanced load and DynamoDB Adaptive Capacity



Poor distribution of traffic across the indexed data

# Dealing with concentrated read throughput



## DynamoDB Accelerator (DAX)

Fully managed, highly available: handles all software management, fault tolerant, replication across multi-AZs within a region

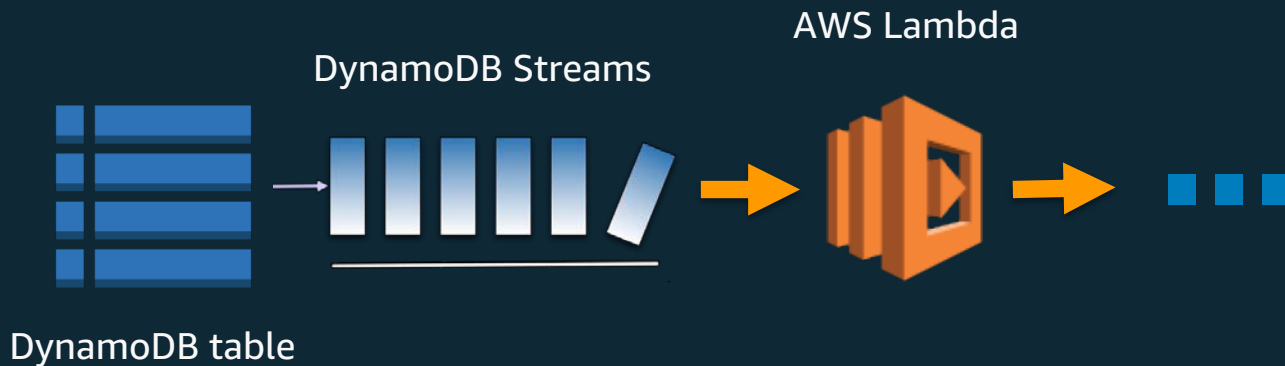
DynamoDB API compatible: seamlessly caches DynamoDB API calls, no application re-writes required

Write-through: DAX handles caching for writes

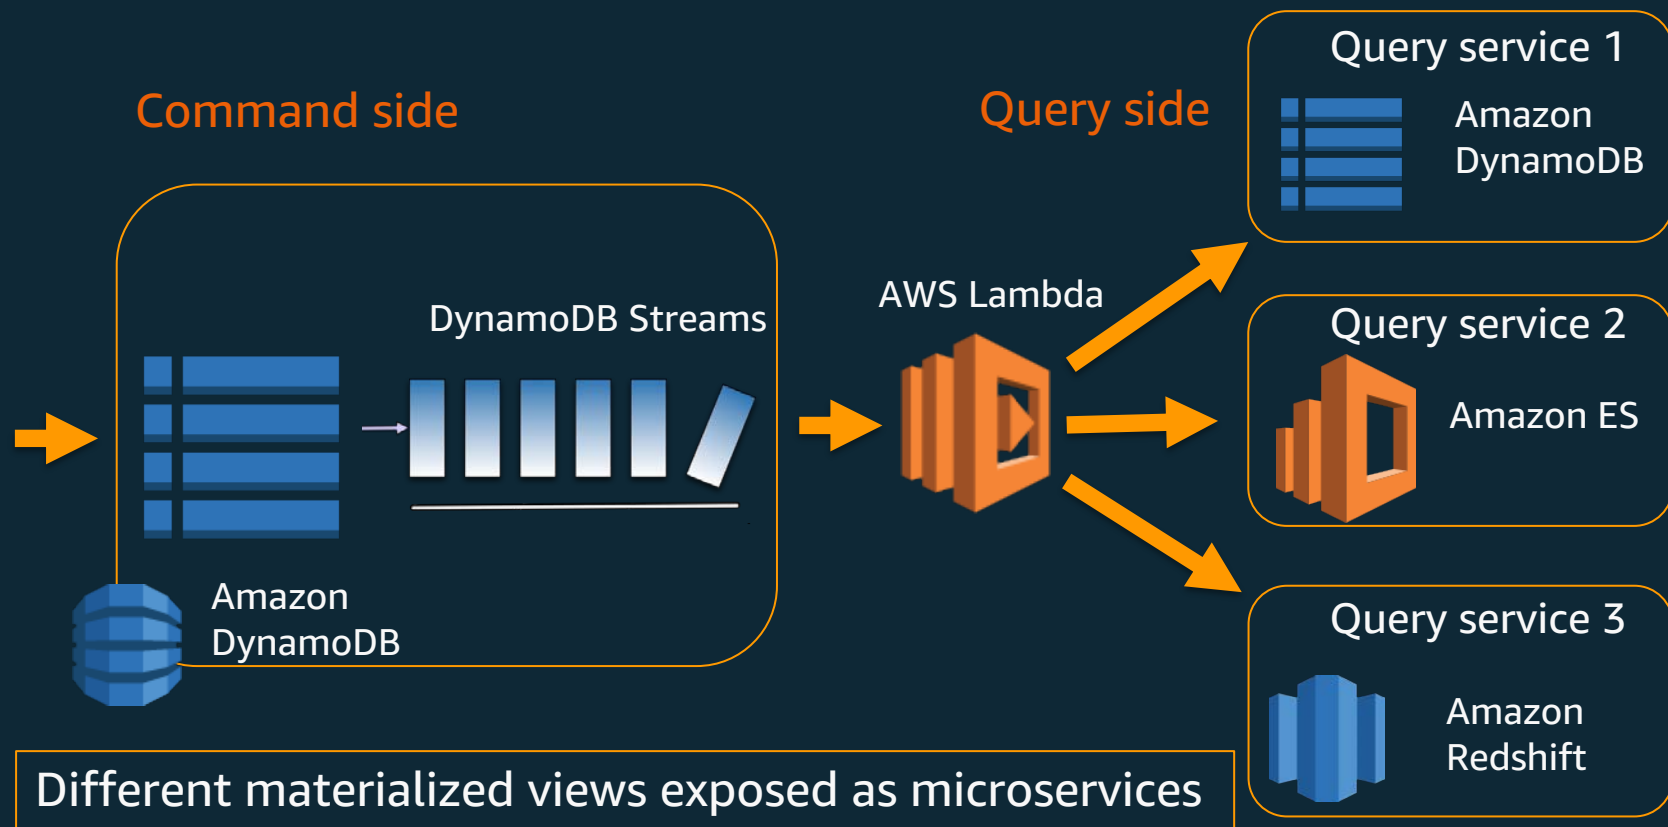
# Integrating DynamoDB into your data flow

# Complex queries and analytics

- Use the service that best meets the requirements
  - Amazon Athena, Amazon Redshift, Amazon Elasticsearch Service...
- Deliver data updates reliably with DynamoDB Streams
  - Integrates with AWS Lambda
  - End-to-end serverless



# Querying in microservices architectures



# Closing summary

- Data management at internet scale gave rise to DynamoDB
  - and to polyglot persistence: use the right database for each job
- DynamoDB: highly-automated (serverless) distributed database
  - ideal for mission-critical OLTP use cases
- Remove scaling concerns – distribute your data and traffic
- You *can* have data consistency and integrity in DynamoDB
- You *can* model relationships beyond key-value in DynamoDB
- Distributed databases are hard to operate - use DynamoDB!



# Thank you!